CS1674: HW10P **Randyll Bearer** **4/12/2018**

First, as a general note about my thoughts on computer vision as a whole, is that I am incredibly impressed. I tried to find dozens of photographs which I thought would “trick” the demos but they were all able to detect the objects within them accurately. From optical illusions to comic book characters the state of the art seemed to perform very well.

For the most part, pictures of animals and humans seemed to be detected accurately the most often. Stationary “Freeze-Frame” photos of normally fast-paced/dynamic objects (e.g. Tornado, Spinning Tennis Ball) seemed to be detected incorrectly the most often. As for the VQA demo in part 3, the demo seemed quite good at distinguishing human features from comic book representations, but wasn’t always good at correctly answering questions of the “How many of X are in this image?” variety, though it got somewhat close. As for what we still need to solve in computer vision based off these findings, we need a way to better detect the number of specific objects of the same class in an image. It would appear that these demos struggle with the overlapping of similar objects.